

Escaping AI slop: a practical guide to writing that sounds human

The fundamental problem with AI-generated text isn't accuracy—it's the absence of a human behind the words. AI predicts the most statistically probable next token; (Undetectable AI) (AI47Labs) human writers make choices based on what they uniquely perceive, feel, and want to communicate. The result: AI writing feels like a Wikipedia article trying to please everyone while offending no one. Your housing justice document likely reads as AI-generated because it exhibits the telltale markers—predictable structures, hedged language, vocabulary nobody actually uses, and emotional flatness even when discussing urgent injustice.

This guide synthesizes research on AI detection, writing craft, and political rhetoric to provide concrete tools for revision. The goal isn't to "trick" readers into thinking a human wrote something; it's to actually write like a human—with voice, conviction, and the rough edges that signal genuine thought.

The technical fingerprint of AI writing

Detection tools identify AI text through two primary measurements: **perplexity** (how predictable word choices are) and **burstiness** (how much sentence length and structure vary). (Twixify) (Detecting AI) AI scores low on both— (Gonzaga) meaning it selects expected words in uniform patterns. Human writing, by contrast, shows higher perplexity (unexpected choices) and burstiness (varied rhythm). (QuillBot)

What this means practically: AI writing lacks surprise. Every sentence lands where you expect it to. Every paragraph follows the same topic-support-conclusion formula. The vocabulary feels slightly elevated but strangely generic—words that sound sophisticated but nobody actually says aloud. As one Berkeley professor noted after thirty years of teaching: words like "delve" and "multifaceted" suddenly appeared everywhere in student writing post-ChatGPT. (The Boston Globe)

The statistical tells include **uniform sentence length** throughout a document, **perfect grammar** with no contractions, (THE DECODER) **long conclusions** that merely restate earlier content, (Pangram Labs) (Substack) and what researchers call the "people-pleasing problem"—prose so cautious about being wrong that it avoids saying anything definite. (Pangram Labs)

Words and phrases that scream AI

Certain vocabulary has become so associated with AI that using it triggers immediate suspicion. These words appeared rarely in human writing before 2023 but now saturate AI output: [Wikipedia](#)

Extreme red flags: delve, tapestry, realm, nestled, meticulous, pivotal, landscape, intricate, kaleidoscope, showcase, underscore, foster, nuanced, multifaceted, embark, navigate, leverage, holistic, robust, seamless, groundbreaking, transformative

Phrases to eliminate entirely: "In today's fast-paced world," "It's worth noting that," "Let's delve into," "rich tapestry of," "navigate the complexities," "unlock the potential," "at the forefront of," "a testament to," "bridging the gap between"

Opening patterns that signal AI: "In the ever-evolving landscape of..." / "Welcome to the world of..." / "Picture this..." / "Ever wondered...?"

Conclusion patterns equally damning: "In conclusion..." / "To summarize..."
[Pangram Labs](#) / "By understanding X, you can Y..." / "The bottom line is..."

The deadly "It's not X, it's Y" formula has become perhaps the single most identifiable AI tell: "It's not about posting more. It's about posting smarter." This sentence structure, along with the rhetorical-question-plus-short-answer pattern ("The best part? It works."), reads as AI writing to anyone paying attention. [Plus](#)

Beyond specific phrases, AI exhibits structural patterns: excessive em-dashes, bullet points where prose would be stronger, paragraphs of suspiciously uniform length, [Pangram Labs](#) and what editors call the "assistant voice"—overly helpful, overly organized, bristling with signposts ("First... Second... Third..."). [Wikipedia](#)

What makes human writing feel human

The research is clear: authenticity comes from **voice, imperfection, specificity, and stakes**—precisely the elements AI training systematically suppresses.

Voice isn't a style; it's a stance. Writer Sonya Huber describes voice as "a chorus of voices" expressing different facets of self—shaped by experience, relationships, community. [Literary Hub](#) AI aggregates so many perspectives it has none itself. [Pangram Labs](#) As AI researcher Nathan Lambert notes: "The cultivation of voice is one of my biggest recommendations... Voice is where you process information into a unique representation—this is often what makes information compelling." Language model providers train for "the average user," optimizing against the quirks that make individual voices distinctive.

Imperfection signals humanity. The absence of typos, hesitations, self-corrections, and rule-breaking reads as uncanny. Human writing breaks grammar rules for effect—[Undetectable AI](#) sentence fragments, starting with "And" or "But," split infinitives. AI's perfect grammar becomes a tell. [Pangram Labs](#) One writer put it sharply: "AI loves polished sentences. Smooth transitions. Predictable rhythms. But perfection is boring. Your readers don't want flawless paragraphs. They want your voice." [Medium](#)

Specificity creates connection. Compare "She fled an unhappy marriage" with "One day she packed a bag and walked out and never looked back, vowing no one would ever make her feel invisible or small again." The second version lets readers fill in details from their own experience. [Writers in the Storm](#) AI produces abstract language that could fit many contexts; [Pangram Labs](#) human writers ground abstractions in sensory, concrete detail.

Stakes generate energy. The best political writing comes from genuine investment—what Sonya Huber calls "deep feeling with a subject." AI writing feels emotionally flat because nothing is actually at stake for the writer. This shows in hedged language, qualified claims, and the careful avoidance of strong positions—all trained behaviors that make AI "safe" but strip writing of urgency. [Pangram Labs](#)

Lessons from essayists who achieve distinctive voice

George Orwell's rules remain foundational. From "Politics and the English Language" (1946): Never use a metaphor you're used to seeing in print. Never use a long word where a short one will do. If it's possible to cut a word out, cut it out. Never use the passive where you can use the active. And most importantly: break any rule sooner than say anything outright barbarous.

Orwell's core warning applies directly to AI-sounding prose: "You can shirk [this trouble] by simply throwing your mind open and letting the ready-made phrases come crowding in. They will construct your sentences for you—even think your thoughts for you." [The Orwell Foundation](#)

Joan Didion understood sentence structure as meaning. "To shift the structure of a sentence alters the meaning of that sentence, as definitely and inflexibly as the position of a camera alters the meaning of the object photographed." [Miller's Book Review](#) She varied sentence length deliberately—for every long sentence, a short one for rhythm. [Substack](#) She withheld information strategically: "Don't dump it all out at once." [Substack](#) Zadie Smith wrote of Didion: "I realized that a woman could speak without hedging her bets, without hemming and hawing, without making nice, without poeticisms, without sounding pleasant or sweet, without deference, and even without doubt." [Gvprssa](#)

Ta-Nehisi Coates emphasizes conviction and revision. "The best writing unabashedly shares a clear opinion." Writing is "10% the idea and 90% continuing to revise, revise, revise." (Nothing Any Good) His work on race came from years of research and lived experience—what he calls doing "actual reporting" behind the rhetoric. "Persistence is underrated and talent is incredibly overrated." (Oprah Daily)

James Baldwin aimed for sentences "clean as bone." "Don't describe it, show it. That's what I try to teach all young writers—take it out! Don't describe a purple sunset, make me see that it is purple." (Literary Hub) For Baldwin, writing was moral: "If there is no moral question, there is no reason to write." (Literary Hub)

The common thread: these writers take positions. They cut ruthlessly. They vary rhythm intentionally. They draw on specific experience. They revise extensively. None of this describes AI output.

What makes political writing compelling rather than preachy

The line between persuasive and preachy is about reader agency. **Preachy writing inflates the author's moral authority while minimizing the reader's capacity for judgment.** It uses "order-heavy words" and expects audiences to follow without engagement. Persuasive writing invites conversation rather than demanding compliance.

Effective political prose makes issues personal through concrete stories. In analyzing successful tenant organizing, one study noted that KC Tenants—one of the most effective tenant unions in the country—leads with specific situations: "The heat was out during one of the coldest weekends... students in the building... \$350 a month for like 8 years... 'we can't go. We won't go.'" That last line—simple, direct, human—does more than abstract demands ever could.

Avoid familiar political terminology. Monica Drake's *The Folly of Loving Life* never uses the word "gentrification" though it's a central theme. "When it comes to hot-button political topics, language that's too familiar can cause us to shut down instead of open up." (litreactor) This applies to housing justice language: terms like "affordable housing crisis" or "housing insecurity" can trigger automatic responses rather than genuine engagement.

Nuance strengthens rather than weakens. Cari Luna's novel about NYC squatters includes "idealists, pragmatists, pessimists, and opportunists... No one is wholly good here, or wholly bad—but that doesn't keep the novel from pointing out that the way these people were shoved aside is a goddamned shame." (litreactor)

Channel anger purposefully. Audre Lorde: "Every woman has a well-stocked arsenal of anger potentially useful against those oppressions... Focused with precision it can become a powerful source of energy serving progress and change." The civil rights movement channeled "sick and tired of being sick and tired" into organized action. ACT UP's anger "wasn't chaotic—it was focused, calculated, and purposeful."

Successful manifestos combine critique with inauguration—they criticize the present while announcing its passing. MLK's "I Have a Dream" connected civil rights to founding American values rather than opposing them. It used specific metaphors ("bad check" for unfulfilled promises) and repetition for rhythm. Brainly

Practical revision checklist

Structural patterns to break:

- Eliminate the introduction-three-points-conclusion format entirely
- Vary paragraph length dramatically—some one sentence, others much longer
- Cut signposting phrases ("First," "Additionally," "In conclusion")
- Remove any conclusion that merely restates earlier points

Language to purge:

- Search for the AI vocabulary list above and replace or delete
- Cut hedging language: "arguably," "to some extent," "it could be said"
- Replace abstract nouns with concrete images
- Use contractions where natural speech would
- Delete any phrase you've seen AI produce

Voice to develop:

- Write at least one sentence per section that only you would write
- Include one specific detail that comes from your actual experience or research
- Take a position without hedging—say what you actually think
- Let some sentences break grammar rules for emphasis
- Read aloud and cut anything that sounds like it's performing intelligence

Rhythm to vary:

- Follow long sentences with short ones. Like this.
- Use fragments deliberately. For emphasis.
- Vary sentence openings—don't start every sentence with subject-verb
- Let some paragraphs breathe; compress others

Political writing specifically:

- Lead with a specific human situation, not an abstract claim
 - Frame demands concretely ("30 percent rent reduction") not abstractly ("affordable housing")
 - Show how abstract injustice lands on specific bodies
 - Channel anger into clear demands rather than diffuse outrage
 - Invite readers into the struggle rather than lecturing them about it
-

The deeper principle

The reason AI writing fails isn't technical—it's philosophical. **AI writes information; humans write experience.** (Medium) AI predicts what comes next based on patterns; human writers make choices based on what they uniquely see, feel, and believe.

(Undetectable AI) Voice emerges from the intersection of personality, experience, and deliberate craft— (Scribophile) (ProWritingAid) exactly what training on aggregate human output systematically erases.

For your housing justice document, this means asking: What do you actually know about housing injustice that a language model doesn't? What specific situations, faces, moments inform your understanding? What position are you willing to take that might make some readers uncomfortable? What rough edges reveal genuine struggle with the material rather than polished detachment?

The goal isn't smuggling human-sounding prose past detection algorithms. It's recovering the human perspective that got lost when good research passed through AI's smoothing function. Somewhere in your document, there's genuine conviction about housing as a human right. Find it. Let it sound like something a person would actually say. That's what readers—and movements—need.